

AI Lab - Session 1

Search Algorithms

Riccardo Sartea

University of Verona
Department of Computer Science

March 15th 2018



UNIVERSITÀ
di **VERONA**

Dipartimento
di **INFORMATICA**

1 Introduction

- OpenAI Gym
- Environment Characteristics
- Installation
- Guidelines

2 Practice

The OpenAI Gym Framework

What is it

Gym is a toolkit for developing and comparing reinforcement learning algorithms. It supports teaching agents everything from walking to playing games like Pong or Pinball

What is it for

- An open-source collection of environments that can be used for benchmarks
- A standardized set of tools to define and to work with environments

Where to find it

<https://gym.openai.com>

Environments for the First Session

- In this first lab session all the environments are grid-worlds
- Every cell has its (x, y) coordinates and a content, e.g., start (S), wall (W), goal (G)...
- A state is identified by an integer
- An action is identified by an integer


Variables:

- *action_space* - the space of possible actions: usually a range of integers $[0, \dots, n]$
- *observation_space* - the space possible observations (states): usually a range of integers $[0, \dots, n]$
- *startstate* - start state (only one)
- *goalstate* - goal state (only one)
- *grid* - flattened grid (1-dimensional array)

Methods:

- *render()* - renders the environment
- *sample(state, action)* - returns a new state sampled from the ones that can be reached from *state* executing *action*
- *pos_to_state(x, y)* - returns the state given its position in x and y coordinates
- *state_to_pos(state)* - returns the coordinates x and y of a state

The SmallMaze Environment



Listing 1: Installation

```
mkdir aicourse
cd aicourse
git clone https://github.com/openai/gym
cd gym
pip3 install --user -e .
cd ..
git clone https://github.com/SaricVr/gym-ai-lab
cd gym-ai-lab
pip3 install --user -l --no-deps -e .
cd ..
git clone https://github.com/SaricVr/ai-lab
```


- You will be working within the **ai-lab** repository
- Each algorithm to implement has its already available prototypes of the functions, e.g., `[search.algorithms.ids(problem, stype)]`
- In order to learn how to use the main features of OpenAI Gym environments you can look at some example code in `[session1/examplecode.py]`
- Different fringe datastructures are already at your disposal `[datastructures.fringe]`. *PriorityFringe* handles priority using *FringeNode.value* property
- Every step, i.e., moving from a state to another, has cost 1. *FringeNode.pathcost* then will always be equal to the depth of such node
- Use the L_1 norm as distance heuristic between a state and the goal `[search.heuristics.l1_norm(p1, p2)]`

- Files [session1/graph_search.py] and [session1/tree_search.py] are the “main” entry point for graph search and tree search algorithms respectively, i.e., the runnable files
- The solution of your search algorithms must be sequences of state identifiers (integers) or *None*
- Along with the solution, you must also return the stats (time, nexp, maxmem) of you algorithms:
 - ▶ Time elapsed between the start and the end of the algorithm
 - ▶ Number of states expanded
 - ▶ Maximum number of states in memory at the same time (fringe + closed)
- Validate carefully the solutions and also the stats returned by your code

1 Introduction

2 Practice

- Uninformed Search
- Informed Search

Iterative Deepening Search (IDS)

```
1: function ITERATIVE-DEEPENING-SEARCH(problem)
2:   for depth  $\leftarrow$  0 to  $\infty$  do
3:     result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
4:     if result  $\neq$  CUTOFF then return result

5: function DEPTH-LIMITED-SEARCH(problem, limit)
6:   return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)

7: function RECURSIVE-DLS(node, problem, limit)
8:   if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
9:   if limit = 0 then return CUTOFF
10:  cutoff_occurred  $\leftarrow$  False
11:  for each action in problem.ACTIONS(node.STATE) do
12:    child  $\leftarrow$  CHILD-NODE(problem, node, action)
13:    result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)
14:    if result = CUTOFF then cutoff_occurred  $\leftarrow$  True
15:    else if result  $\neq$  FAILURE then return result
16:  if cutoff_occurred then return CUTOFF
17:  return FAILURE
```

Note: this is a **tree search** version

Task 1

Implement the **IDS (iterative deepening search)** algorithm based on **tree search**. Work with the *SmallMaze-v0* environment. Proceed by steps:

- 1 Implement the **search.algorithms.dls_ts(problem, limit)** procedure. Test it and debug it with different limit values
- 2 Implement the **search.algorithms.ids(problem, stype)** procedure based on **search.algorithms.dls_ts(problem, limit)**

Task 2

Implement the **IDS (iterative deepening search)** algorithm based on **graph search** [`search.algorithms.dls_gs(problem, limit)`] version. Work with the *SmallMaze-v0* environment.

Are there any differences in results between task 1 and task 2? If so, why?

Uniform-Cost Search (UCS)

Input: *problem*

Output: *solution*

```
1: node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
2: fringe  $\leftarrow$  PRIORITY-QUEUE ordered by PATH-COST, with node as the only element
3: closed  $\leftarrow \emptyset$ 
4: loop
5:   if IS-EMPTY(fringe) then return FAILURE
6:   node  $\leftarrow$  REMOVE(fringe) ▷ Remove node with highest priority
7:   if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
8:   closed  $\leftarrow$  closed  $\cup$  node
9:   for each action in problem.ACTIONS(node.STATE) do
10:     child  $\leftarrow$  CHILD-NODE(problem, node, action)
11:     if child.STATE not in fringe and child.STATE not in closed then
12:       fringe  $\leftarrow$  INSERT(child, fringe)
13:     else if child.STATE in fringe with higher PATH-COST then
14:       replace that fringe node with child
```

Note: this is a **graph search** version

Implement the **UCS (uniform-cost search)** algorithm based on **graph search**. Work with the *SmallMaze-v0* environment. Proceed by steps:

- 1 Implement the general `search.algorithms.graph_search(problem, fringe, f)` procedure
- 2 Verify that `search.algorithms.ucs(problem, stype)` with *stype = search.algorithms.graph_search* works as expected

Task 4

Implement the **UCS (uniform-cost search)** algorithm based on **tree search**. Work with the *SmallMaze-v0* environment. Proceed by steps:

- 1 Implement the general `search.algorithms.tree_search(problem, fringe, f)` procedure
- 2 Verify that `search.algorithms.ucs(problem, stype)` with `stype = search.algorithms.tree_search` works as expected

Are there any differences in results between task 3 and task 4? If so, why?

Task 5

Analyze and run the code of the BFS¹ (breadth-first search) algorithm in both graph search and tree search versions. Work with the *SmallMaze-v0* environment.

What can you say about the results?

¹Remember that the BFS does not use any evaluation function

Task 6

Implement the following algorithms based on `search.algorithms.graph_search(problem, fringe, f)` and `search.algorithms.tree_search(problem, fringe, f)`:

- Greedy best-first [`search.algorithms.greedy(problem, stype)`]
- A* best-first [`search.algorithms.astar(problem, stype)`]

Work with the *SmallMaze-v0* environment.

hint: Have a look at the implementation of UCS in `algorithms.py` to see how you can define and use a cost function.

Modify this to be the heuristic for Greedy and the heuristic + cost for A*. You can find already implemented heuristics in `heuristics.py`

Task 7

Run all of your your search algorithms (graph and tree based versions) on the *SmallMaze-v0*, *GrdMaze-v0* and *BlockedMaze-v0*. Discuss the results: which are optimal? Which are not and why?

Compare your results

You can compare the results achieved by your solution with the one provided by us.

Have a look at the following files:

- Results for tree search
- Results for graph search

Alternatively you can update the git repository and then look at the files

`session1/graph_search_results.txt` and
`session1/tree_search_results.txt`

To update the repository do the following:

```
cd ai-lab
git commit -a -m "a message"
git pull
```